Predicting The Unpredictable

- Leveraging Deep Learning for Counterterrorism -

Oluwasegun Adegoke
George Mason University
Fairfax, Virginia
oadegoke@gmu.edu

Abstract

This project delves into the predictive modeling of terrorist activities using advanced deep learning techniques. Leveraging a comprehensive dataset spanning several years, the research constructs and compares time-series models based on Long Short-Term Memory (LSTM) networks and Graph Convolutional Networks (GCN) to forecast the frequency and type of future attacks. The results showcase the effectiveness of LSTM in capturing temporal dependencies and GCN in unraveling the complex, interwoven patterns of the terrorist network structure. This paper contributes to the field by improving predictive accuracies, offering insights into terrorist behaviors, and presenting novel applications of deep learning in counterterrorism efforts.

Keywords: Deep Learning, LSTM, GCN, Terrorism Prediction, Network Analysis, Predictive Analytics, Counterterrorism Intelligence

I. INTRODUCTION

The ubiquity of global terrorism poses significant challenges to national security and public safety. Understanding and forecasting terrorist activities enable proactive countermeasures, making predictive analytics an indispensable tool for security agencies. Historical terrorism data embodies complex patterns of spatial and temporal nature, demanding sophisticated analytical approaches for effective interpretation. This research harnesses the potential of deep learning, employing LSTM and GCN models to discern and predict terrorist activity patterns.

The utilization of LSTM allows for the examination of temporal sequences and the identification of underlying patterns in the occurrence of terrorist events. Concurrently, GCN's incorporation enables the exploration of the intricate relationships between various entities within the terrorism context, such as locations, attack types, and dates. The models' predictive capabilities aim to assist in preemptive planning and allocation of resources to mitigate potential threats.

The succeeding sections will detail the methods employed in data preprocessing, feature extraction, and model architecture, followed by a presentation of the training procedures and results. This paper endeavors to bridge the gap between deep learning and counterterrorism, extending the frontier of predictive analytics in the domain.

II. Literature Review

Terrorism forecasting and analysis have long been subjects of significant research interest due to their implications for national security and global stability. The existing body of literature reveals a multidisciplinary approach, integrating political science, computer science, and statistics. Early works primarily focused on statistical methods for understanding and predicting terrorist activities, with Clauset, Young, and Gleditsch (2007) applying power-law distributions to the severity and frequency of terrorist attacks.

The advent of machine learning and deep learning has shifted the focus towards data-driven models. LSTM networks, first introduced by Hochreiter and Schmidhuber (1997), have been pivotal in analyzing time-series data for terrorism patterns due to their ability to capture long-term dependencies (Liang et al., 2016). Recently, researchers have extended traditional analyses by implementing Graph Neural Networks (GNNs) to elucidate the complex web of relationships within terrorist networks (Christopher et al, 2015). While the use of GNNs in this domain is nascent,

initial studies indicate their potential for uncovering non-obvious insights about the operational structures and connections between different entities involved in terrorism.

III. Dataset Overview

The primary source of data for this project is the Global Terrorism Database (GTD), which is an open-source database including information on terrorist events around the world from 1970 through 2016. This database is maintained by the National Consortium for the Study of Terrorism and Responses to Terrorism (START), based at the University of Maryland. It provides comprehensive data on domestic as well as international terrorist incidents that have occurred during this time frame and includes systematic data on the date and location of incidents, the weapons used, nature of the target, the number of casualties, and the group or individual responsible.

This dataset contains information, initially including 170,350 rows and 135 columns. After cleaning and selection processes, the data was refined to a more manageable size of 165,744 rows and 22 columns. This preprocessed data incorporates a mix of categorical, numerical, and textual data types. This comprehensive range of data types allows for a detailed examination of the specific characteristics of each terrorist incident.

Key Features

The primary features used in our analysis include:

- **Date of Event**: The exact date when the terrorist attack occurred. This feature is crucial for analyzing temporal patterns and trends in terrorism.
- **Location Information**: Includes country, region, and exact latitude and longitude. These geographical details are vital for studying the spatial distribution of terrorist activities.
- **Attack Type**: Describes the method of attack (e.g., bombing/explosion, armed assault, hijacking). This categorical feature is used as the target variable in our predictive model.
- **Weapon Type**: Specifies the type of weapon used in the attack, providing insights into the operational aspects of terrorist incidents.
- **Target Type**: Details the type of target/victim (e.g., private citizens, police, military), which helps in understanding the preferences or biases of terrorist groups.
- **Casualty data** which includes numbers on deaths, injuries, and overall casualties.
- **Additional context** such as whether the attack occurred on a holiday, and days since the last attack.

Preprocessing and Handling Missing Data

Given the comprehensive nature of the GTD, several features had missing or incomplete entries. Handling these missing values was a critical step in preparing the dataset for analysis:

- **Mode Imputation**: For categorical data such as 'Attack Type' and 'Target Type', where missing values were less frequent, the mode (most common category) of each feature was used to fill in gaps. This approach was chosen to maintain the statistical integrity of the dataset.
- **Median Imputation**: For continuous variables like latitude and longitude, which displayed skewed distributions, median values were used for imputation to avoid the influence of outliers.
- **Deletion**: In cases where a significant portion of crucial data was missing (more than 30% missing data), such as detailed descriptions of the perpetrators, those records were excluded from the analysis to ensure the robustness and accuracy of the model predictions.

Ethical Considerations

In handling data related to terrorism, ethical considerations were meticulously observed. The GTD itself is anonymized, with no information on individuals or sensitive personal data. This project focuses solely on aggregated trends and patterns, avoiding any analysis that could lead to discrimination or harm against specific groups or individuals.

IV. Methodology

Time-Series Analysis Using LSTM Networks

Long Short-Term Memory networks (LSTMs) are a class of Recurrent Neural Networks (RNNs) that are capable of learning order dependence in sequence prediction problems. This project employs LSTM networks to analyze time-series data related to terrorist activities. The rationale for selecting LSTMs lies in their proven track record of handling vanishing gradient problems and capturing temporal dependencies in sequential data, which is essential for accurate predictions of the timing and frequency of terrorist events.

Implementation of Graph Neural Networks (GNNs)

The complex interactions and dependencies between various entities involved in terrorism are modeled using Graph Neural Networks (GNNs), particularly Graph Convolutional Networks (GCNs). GCNs are a powerful neural network architecture for machine learning on graphs. They operate by learning a function of signals/features on a graph G (represented by nodes and edges) that takes the attributes of its nodes (and potentially the structure of the graph) into account. This model applies GCNs to perform node classification, considering nodes as various entities such as dates, locations, and types of terrorist attacks. The connectivity between the nodes represents relationships, such as the occurrence of attacks across different locations or dates.

By modeling the problem using GNNs, this project extends beyond the capabilities of traditional machine learning models to include relational inductive biases, thereby providing a more nuanced understanding of the patterns and connections inherent in the data.

V. Exploratory Data Analysis

The Exploratory Data Analysis (EDA) section of the project provides a critical examination of the dataset to uncover patterns, anomalies, and correlations within the features related to global terrorist activities.

The "Tree map of Casualties by Country" visual reveals a stark contrast in the impact of terrorism across different nations. Dominant areas for Iraq and Afghanistan within the tree map indicate a significant number of casualties, portraying the intense concentration of terrorist activities in these countries. Smaller areas for countries such as Spain and Peru suggest a relatively lower frequency of such violent incidents, showcasing a wide disparity in the impact of terrorism worldwide.

Figure 1 Casualties by Country

The "Global Map of Terrorist Attacks" illustrates the geographical spread and severity of terrorist incidents. The clusters of blue dots, representing attacks with fewer than 75 casualties, contrast sharply with the red dots, highlighting locations with deadly attacks causing more than 75 casualties.

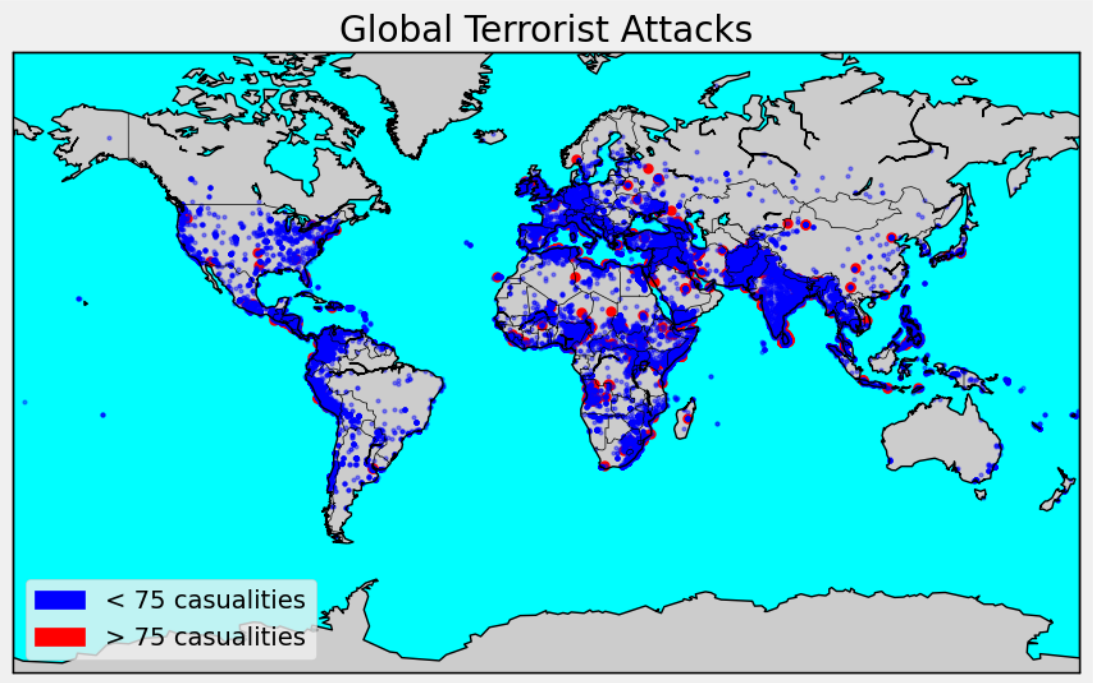

Figure 2 Global Map of Terrorist Attacks

In the "Distribution of Causalities by Region" graph, the disparities in the number of casualties across various global regions are evident. The Middle East & North Africa and South Asia regions stand out with their higher number of casualties, punctuated by outlier points that signify extremely lethal terrorist events.

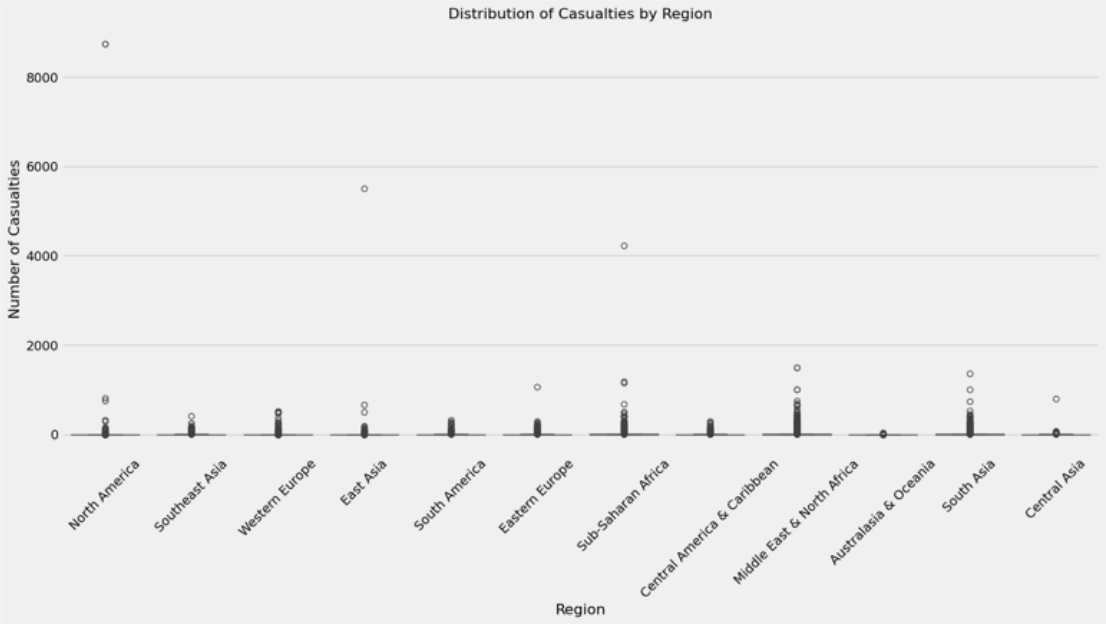

Figure 3 Distribution of Causalities by Region

The "Heatmap of Feature Correlations" provides a succinct and quantifiable view of the relationships between various data features. Strong correlations observed between 'Killed' and 'casualties', as well as 'Wounded' and 'casualties', affirm the significance of these variables as potential predictors for the models.

Finally, the "Number of Terrorist Attacks Per Year" bar chart chronicles the escalation in terrorist activities over the years, with a notable increase in recent times. This upward trend underscores the necessity for predictive modeling in counter-terrorism efforts.

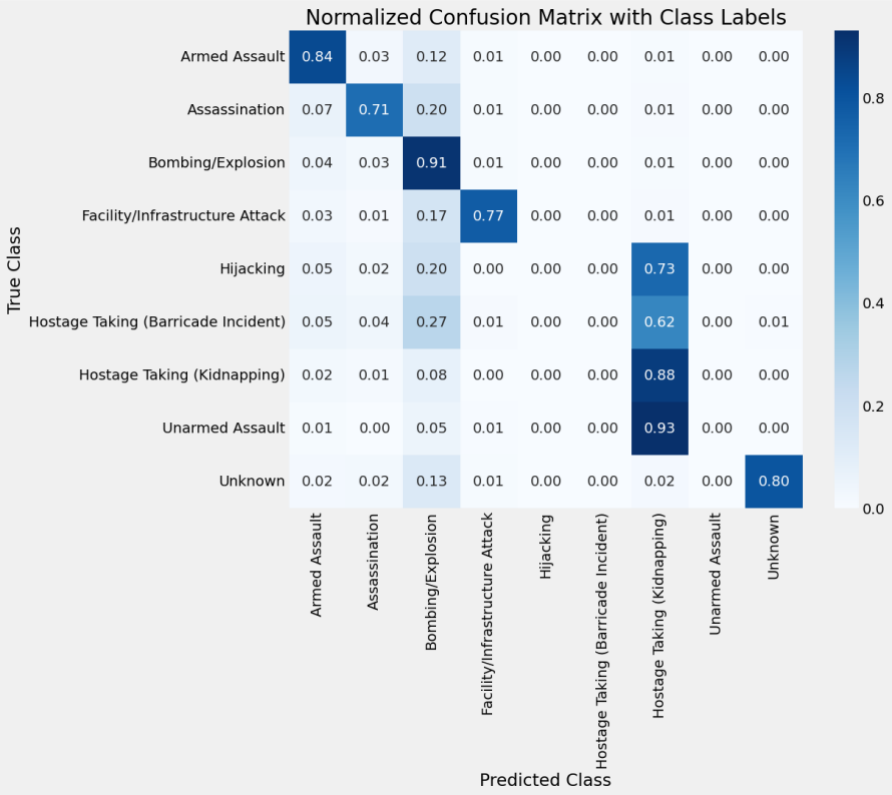

Figure 4 Heatmap of Feature Correlations

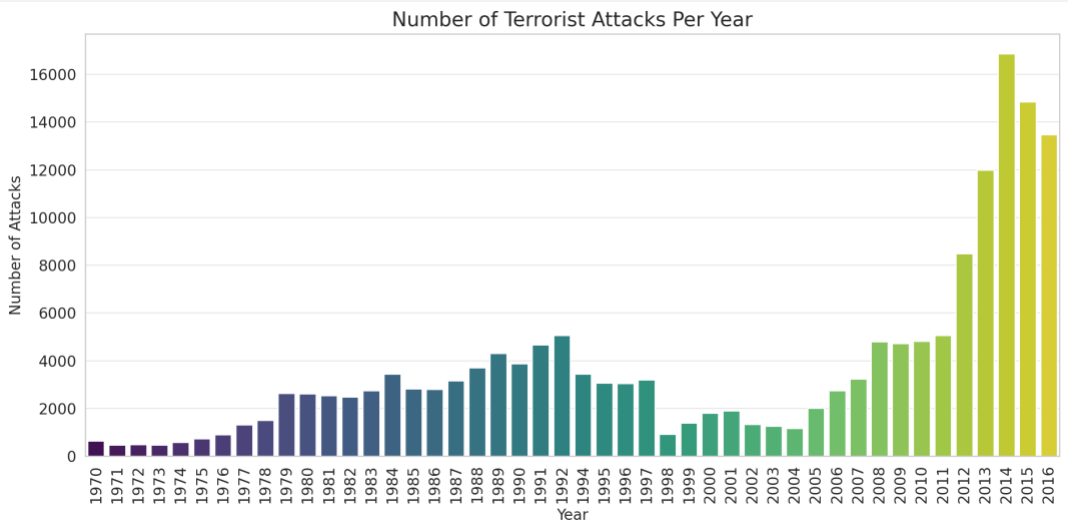

Figure 5 Number of Terrorist Attacks Per Year

VI. Model Architecture and Training

LSTM for Time-Series Analysis

i. LSTM Architecture

The initial approach to forecasting the frequency of terrorist attacks was based on an LSTM network. LSTMs are a type of recurrent neural network (RNN) capable of capturing temporal dependencies in sequence data, making them suitable for time series forecasting. The LSTM model was constructed with an input layer that specifies the shape of the training dataset, followed by two LSTM layers, each with 100 units. The "*ReLU*"activation function was employed to introduce non-linearity to the model, aiding in the capture of complex patterns.

To combat overfitting, dropout layers were inserted after each LSTM layer with a rate of 0.4. The model concludes with a dense layer with a linear activation function to output the forecasted number of attacks. The optimizer used was Adam with a reduced learning rate of 0.0001 to ensure stable and gradual learning, and the model was compiled with a mean squared error (MSE) loss function.

The training process involved feeding the model with monthly aggregated data for 100 epochs, with a batch size of 16, and a validation split of 10%. Early stopping was not utilized in this training, leading to potential overfitting as indicated by divergence between the training and validation loss.

ii. GRU (Gated Recurrent Unit) Architecture

Given the potential overfitting observed with the LSTM model, a second architecture was introduced using Gated Recurrent Units (GRU). GRUs, like LSTMs, are designed to capture dependencies in sequence data but with a simpler structure that can offer computational efficiency.

The GRU architecture mirrored the LSTM's, with two GRU layers of 100 units each and *"ReLU"* activation. Dropout was included with a rate of 0.2 to prevent overfitting. However, this model added regularization to the final dense layer to penalize complex models through L1 and L2 regularization, thus encouraging a simpler model that could generalize better.

The training regime for the GRU model was akin to the LSTM's, except for employing early stopping with a patience of 10 epochs. This means training would cease if the validation loss did not improve for 10 consecutive epochs, with the best model weights being restored from the epoch with the lowest validation loss.

iii. Models Evaluation

Both models were evaluated on a held-out test set. The LSTM model exhibited a high-test loss, indicating that the model may not have generalized well to unseen data. Conversely, the GRU model, with regularization and early stopping, achieved a lower test loss, suggesting a better performance.

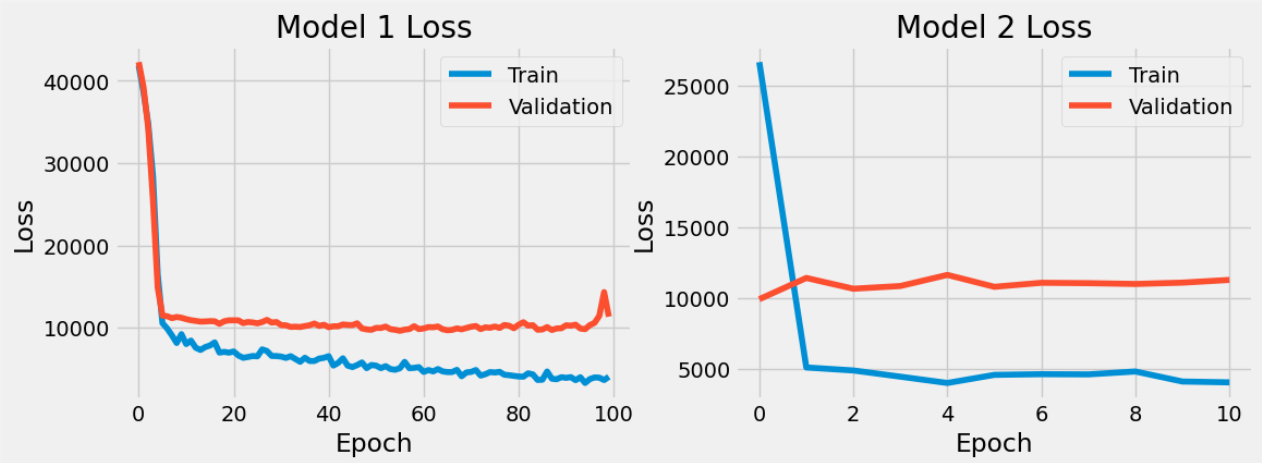

Figure 6 Monthly Sequence Loss Comparison: LSTM vs GRU.

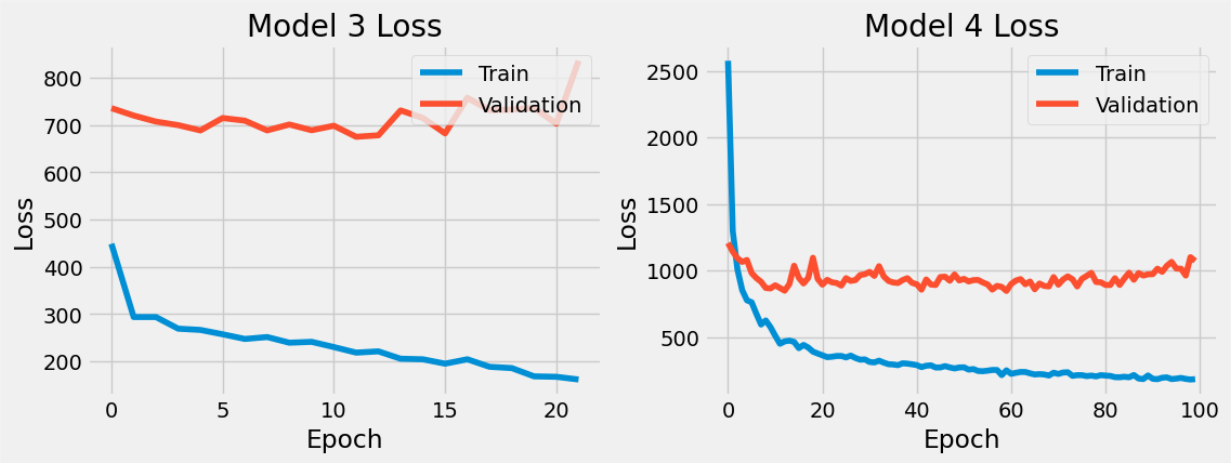

Figure 7 Weekly Sequence Loss Comparison: LSTM vs GRU.

In the evaluation of the models, the LSTM exhibited a steep decline in training and validation loss during the initial epochs, indicating a rapid learning phase. However, the subsequent flattening of the loss curve suggests a plateau in learning, where additional epochs make negligible improvements. This behavior may imply that the LSTM quickly captured the simpler patterns within the dataset but struggled to further refine its predictions to capture more complex patterns, as reflected by the considerable validation loss.

The GRU model, in contrast, showed a more gradual decline in loss, maintaining a consistent improvement across epochs. The closer convergence of the training and validation loss suggests a better generalization to the unseen data. Despite the higher initial loss compared to the LSTM, the GRU's performance improved steadily, without the abrupt changes exhibited by the LSTM, which is indicative of more stable learning.

The models' loss graphs illustrate that while the LSTM was quicker in initial adaptation, the GRU proved more effective over time at understanding and predicting the patterns in the data. This suggests the GRU may be more suited to this task, where complex temporal dynamics are prevalent, and stable learning is necessary for generalization. These findings reinforce the importance of model selection in time-series forecasting of complex phenomena like terrorist activities, where different architectures can lead to markedly different performance outcomes.

Graph Convolutional Networks

i. Graph Representation of Data

To leverage the Graph Convolutional Network (GCN) approach, the dataset must be transformed into a graph format. In this format, nodes represent entities such as locations, dates, and types of attacks, while edges indicate relationships between these entities. In this case, nodes were constructed for each unique location and attack type, with edges connecting locations to corresponding attack incidents. This graph structure encapsulates the complex interplay between different features, enabling the GCN to exploit the rich relational information within the terrorism data.

ii. Graph Convolutional Network Architecture

GCN employed in this project consisted of multiple layers of graph convolution operations that successively aggregate information from a node's neighbors. The GCN aimed to learn a function of features on a graph, which considers the structure of the graph and the features of the nodes and edges. The GCN model used in this research included three layers of graph convolution:

a) The first graph convolution layer takes the initial node features and the adjacency matrix to apply a transformation and aggregation operation, which captures the local neighborhood's structure.
b) The second layer further processes this aggregated information, allowing the model to consider larger neighborhoods.

c) The third and final layer reduces the dimensionality to the number of classes for the multi-class classification task, using a SoftMax activation function to output the probabilities for each class.

Regularization techniques, such as dropout with a rate of 0.5, were utilized after the convolutional operations to mitigate overfitting and enhance the model's ability to generalize to new data. Additionally, the ReLU activation function was employed between layers to introduce non-linearities into the model, crucial for capturing complex patterns.

The model was then trained on the preprocessed and feature-engineered data, including one-hot encoded categorical attributes, and normalized numerical variables. The model was trained for 200 epochs using the Adam optimizer with a learning rate of 0.01. To guide the learning process, a cross-entropy loss function was used due to the multi-class nature of the problem. Training utilized a mask to ensure only the designated training nodes were considered during the optimization process, with the test mask reserved for the final evaluation.

VII. Results

i. Time Series Models

The LSTM and GRU models' predictive capabilities were assessed using Mean Absolute Error (MAE), with separate evaluations for monthly and weekly predictions. The models were trained on historical data and tasked to forecast future attack frequencies, benchmarked against actual observed values.

Table 1 Mean Absolute Error for LSTM VS GRU Models

Mean Absolute Error for LSTM VS GRU		
MODELS	Monthly Sequence	Weekly Sequence
LSTM	236.2774	102.2096
GRU	158.4214	67.6300

The performance of the models on a monthly sequence basis is depicted in *Figure 8*, where the actual vs. predicted attack frequencies for both LSTM and GRU models are visualized. The GRU model consistently outperforms the LSTM, as evidenced by the closer alignment of its predictions with the actual data and a lower MAE value.

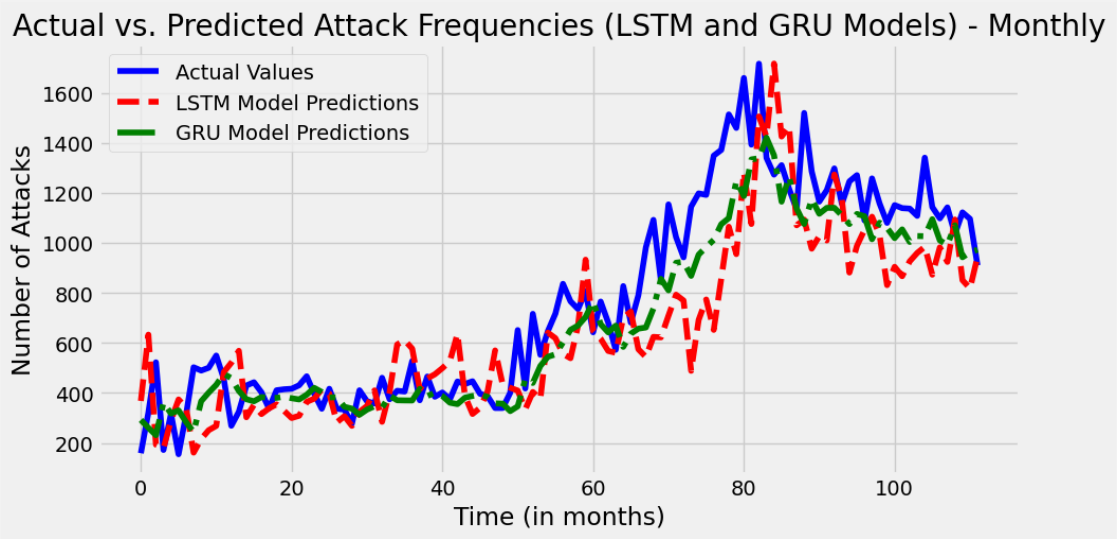

Figure 8 Actual vs. Predicted Attack Frequencies (LSTM and GRU Models) - Monthly

When analyzing weekly sequences, as shown in *Figure 9*, the GRU model demonstrates a more robust predictive accuracy compared to the LSTM. Although both models capture the general trend of attack frequencies, the GRU model's predictions are notably closer to the actual values.

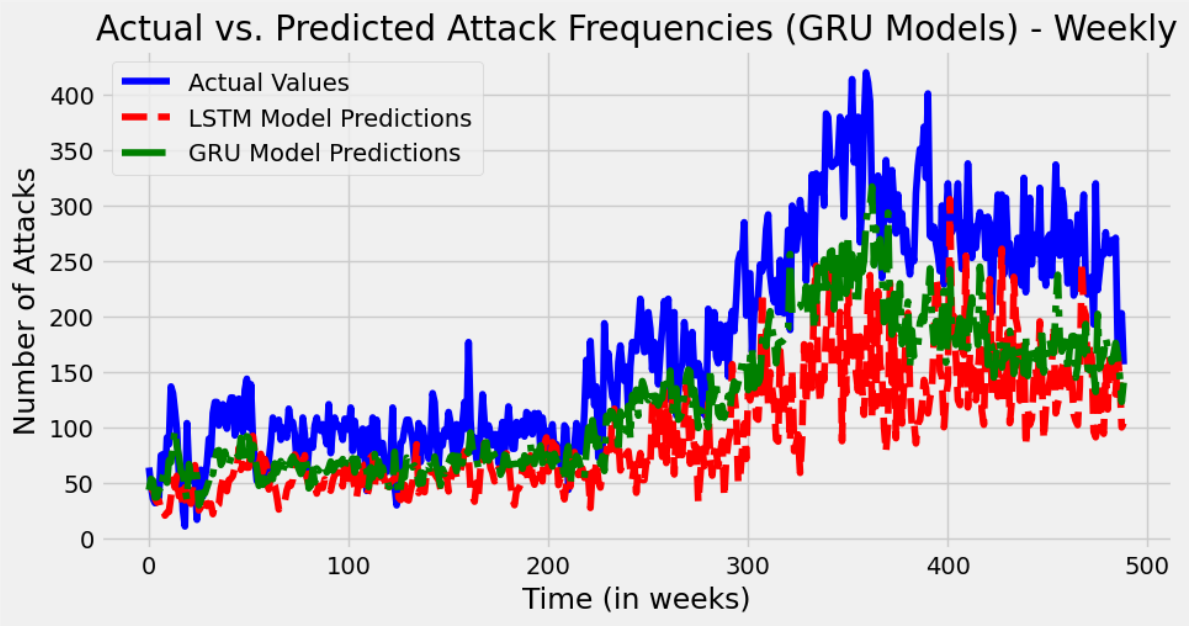

Figure 9 Actual vs. Predicted Attack Frequencies (LSTM and GRU Models) - Weekly

ii. Graph Convolutional Network

The application of GCN has yielded insightful results, revealing the relational structures and influential nodes within the terrorism network. The model achieved an accuracy of 85.01%, with precision, recall, and F1-score all aligning closely, indicating balanced performance across different attack types.

The model's accuracy over epochs is visualized in Figure 10, reflecting a rapid climb to a high accuracy level which sustains over time, surpassing the baseline.

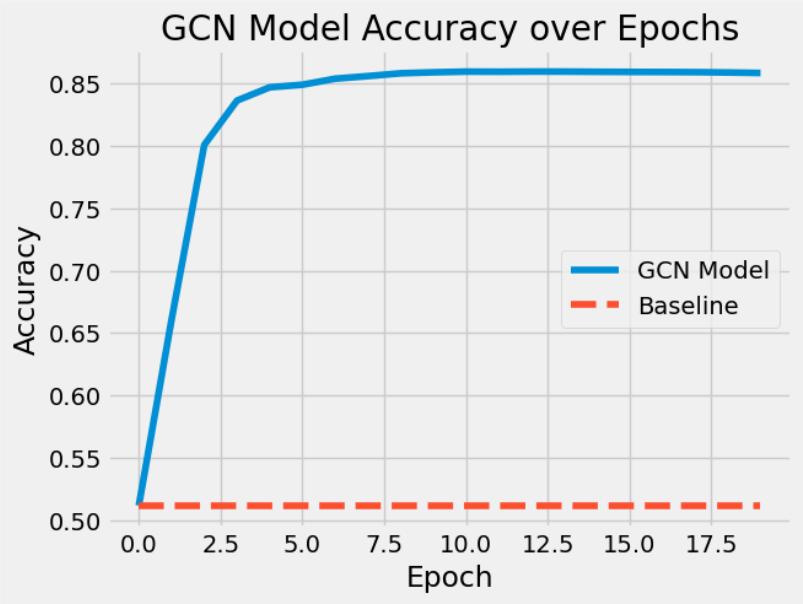

Figure 10 GCN Model Accuracy Over Epochs

Normalized Confusion Matrix, represented in Figure 11, outlines the model's performance across different attack types. It demonstrates high true positive rates for categories like Bombing/Explosion and Facility/Infrastructure Attack, highlighting the model's effectiveness in identifying these attack patterns.

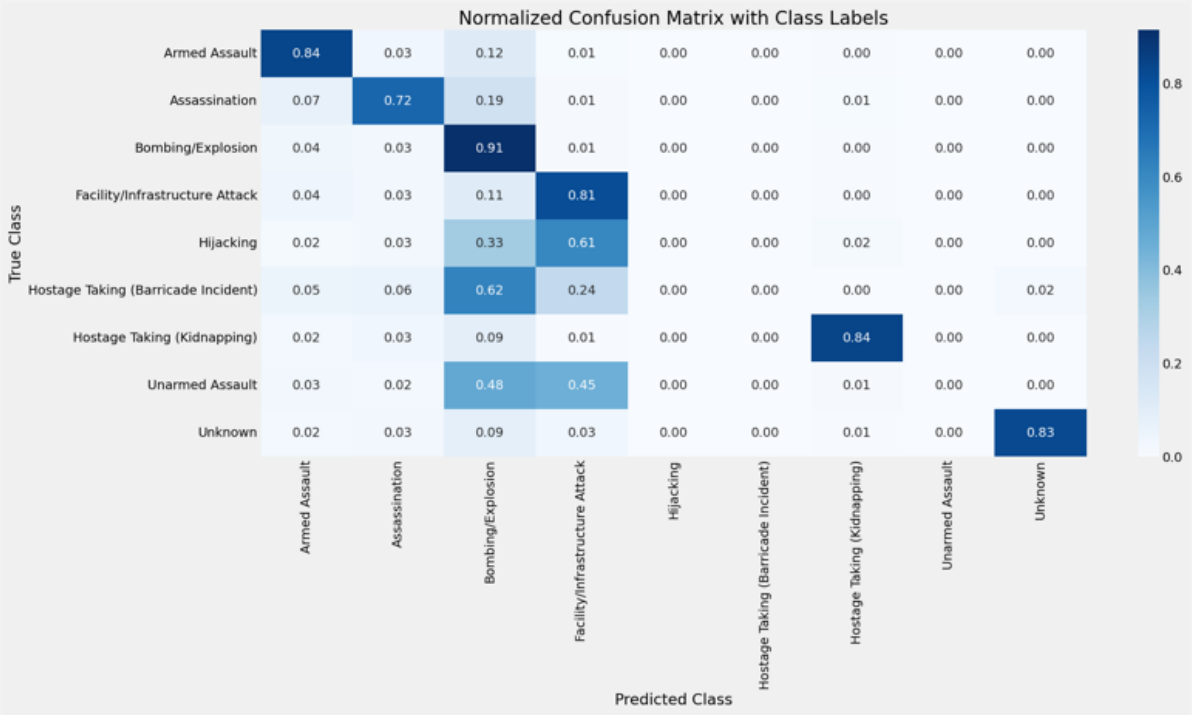

Figure 11 Confusion Matrix with Class Labels

The 'Geographic Distribution of Predicted Attack Types' provides a visual interpretation of the model's ability to geographically discern the type of attacks, emphasizing the relevance of GCN in spatial pattern recognition. Figure 12 conveys how the model distinguishes between attack types across global coordinates, reinforcing the model's spatial predictive capabilities.

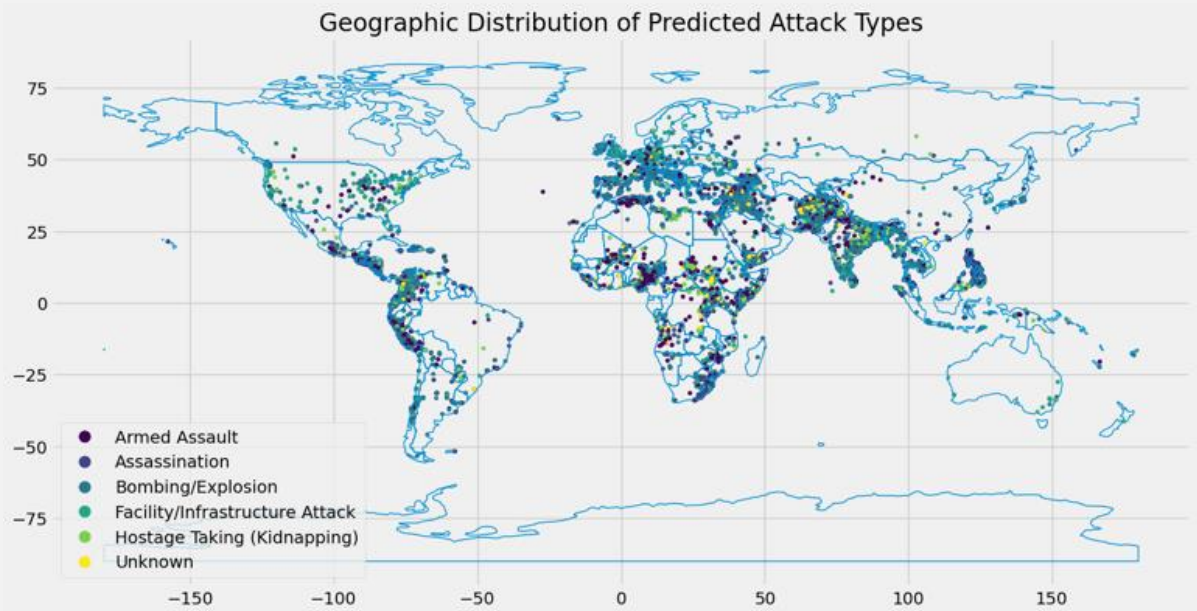

Figure 12 Geographic Distribution of Predicted Attack Types

The identification of the top 10 most influential nodes within the network underpins the model's utility in pinpointing critical elements. For instance, the node 'Bombing/Explosion' holds the highest influence score, suggesting a prevalent attack type with significant impact across the network. This insight, coupled with the identification of critical locations such as 'Baghdad' in Iraq, provides a profound understanding of the network's topology and its centrality in the terrorist landscape.

Top 10 most influential nodes	
Bombing/Explosion	0.309156
Armed Assault	0.289629
Assassination:	0.119341
Hostage Taking (Kidnapping)	0.098470
Facility/Infrastructure Attack	0.085052
('Iraq', 'Middle East & North Africa', 'Baghdad')	0.059966
('United Kingdom', 'Western Europe', 'Belfast')	0.032864
('Pakistan', 'South Asia', 'Karachi')	0.027020
('Iraq', 'Middle East & North Africa', 'Mosul')	0.023151
('Somalia', 'Sub-Saharan Africa', 'Mogadishu')	0.019710

VIII. Discussion

The empirical investigation delineated in this project underscores the potency of deep learning methodologies for prognosticating terrorist incidents. The deployment of LSTM models has elucidated the temporal proclivities of terrorist activities, whilst Graph Convolutional Networks (GCNs) have demystified the intricate interconnections within the terrorist networks. The LSTM model was proficient in initial adaptation to the temporal data; however, it was the GRU model that manifested superior performance over an extensive span, suggesting its heightened efficacy in discerning the elaborate temporal dynamics inherent in the dataset. In contrast, the GCN's ability to discern the relational intricacies within the data and the subsequent prediction of attack patterns signifies an advancement in the identification of high-risk nodes within networks, which could be pivotal in pre-emptive counterterrorism measures.

IX. Conclusion

The conclusive findings of this project pivotally contribute to the realm of counterterrorism analytics. By intricately analyzing historical patterns and network relationships, the LSTM and GRU models illustrate robust capabilities in forecasting the frequency and types of terrorist attacks. The granularity of prediction achieved by these models affords a heightened situational awareness that is vital for pre-emptive security strategizing and resource allocation. Furthermore, the observed predictive precision of the GRU model indicates its applicability in scenarios where long-term dependencies are critical, positing it as a viable instrument for real-world applications. Conversely, the GCN model has emerged as a pivotal tool for elucidating the relational dynamics within terrorist networks, thereby illuminating key nodes and patterns that may otherwise remain concealed.

X. Glossary

- **Deep Learning**: A subset of machine learning involving neural networks with multiple layers that can learn and make intelligent decisions on their own.
- **LSTM (Long Short-Term Memory)**: A type of recurrent neural network that is capable of learning long-term dependencies in sequence data.
- **GRU (Gated Recurrent Unit)**: A streamlined version of LSTM designed to solve the vanishing gradient problem with fewer parameters.

- **GCN (Graph Convolutional Network)**: A neural network architecture designed to process data structured as graphs, capturing the relationships between nodes.
- **Epoch**: An iteration over the entire data set used to train a neural network.
- **Node**: A fundamental unit in a graph, representing entities such as locations or events in the context of GCNs.
- **Temporal Dynamics**: Patterns or trends that occur over time within a dataset.
- **Predictive Analytics**: The use of data, statistical algorithms, and machine learning techniques to identify the likelihood of future outcomes based on historical data.
- **Overfitting**: A modeling error that occurs when a function is too closely fitted to a limited set of data points and may fail to generalize to new data.
- **Regularization**: Techniques used in machine learning to reduce the error by fitting a function appropriately on the given training set and to avoid overfitting.

XI. REFERENCES

1. Clauset, A., Young, M., & Gleditsch, K. S. (2007). On the frequency of severe terrorist events. *Journal of Conflict Resolution, 51*(1), 58-87. https://doi.org/10.1177/0022002706296157
2. Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation, 9*(8), 1735-1780. https://doi.org/10.1162/neco.1997.9.8.1735
3. Bandara, Kasun & Bergmeir, Christoph & Hewamalage, Hansika. (2020). LSTM-MSNet: Leveraging Forecasts on Sets of Related Time Series With Multiple Seasonal Patterns. IEEE Transactions on Neural Networks and Learning Systems. PP. 1-14. 10.1109/TNNLS.2020.2985720.
4. Squicciarini, A. & Rajtmajer, S. & Liu, Y. & Griffin, Christopher. (2015). Identification and characterization of cyberbullying dynamics in an online social network. 280-285. 10.1145/2808797.2809398.